

ADR-0002 — Process, risk, and governed autonomy

Purpose

ADR-0001 settled *how a thing comes to be* in the Lab: human material is digested into content-addressed LogLine Acts. This record settles the complementary half: **how a thing is treated, how it moves, and how the Lab grows — proportionally and safely**, without inventing a metaphysics of who is allowed to be right.

Relationship to ADR-0001

ADR-0001 what a thing IS and how it COMES TO BE (digestion, Act, envelope, custody)
ADR-0002 how a thing is TREATED and MOVES, and how the Lab GROWS (process, risk, automation)
ADR-0003 how a thing is SEEN and what earns ATTENTION (projections v4, attention v5)

Law 1 — Process, not source privilege

The Lab does not treat a statement as true because a human said it, nor false because an LLM said it. It does not treat a file as canonical because it exists, nor a JSON object as safe because it validates. **Every input is material.** The process determines treatment.

Origin does not decide truth.
Form does not decide safety.
Process decides treatment.
Risk decides rite.
Consequence decides weight.

This dissolves the need for any special status for “what the LLM said” vs “what the human said.” Both are participants in processes — with different roles, grants, and confirmation powers — never inherently true. The same human sentence may move through a light process or a heavy one depending entirely on what it would *do*.

The risk matrix (L0–L5) — consequence

Risk classifies the **consequence of one act**. The canonical matrix lives in docs/DANGER_TIERS.md:

L0 read / projection (no external effect)
L1 local ledger write
L2 deterministic bounded local action
L3 external integration, no irreversible effect
L4 external compensable effect (grant + signoff + budget + sandbox)
L5 external irreversible effect (strictest gate)

Risk is a property of the *act's consequence*, not of its origin or its prose.

The scale law

fixed canonical form -> everything is written in one mold (ADR-0001)
+ processual classification -> each thing knows which rite it passes through (Law 1)
+ proportional risk -> each rite carries weight matched to consequence (L0-L5)
= scale

Fixed form prevents chaos; process prevents arbitrariness; proportional risk prevents both recklessness and paralysis. **The Lab scales because it does not decide everything by hand.**

Automation — continuity capped by consequence

Automation is **permission to continue under contract** — how far a process may move without fresh human input. It is a *second axis*, distinct from danger:

L-axis (danger) = consequence of one act (docs/DANGER_TIERS.md)
A-axis (automation) = continuity / self-movement (docs/AUTOMATION_DOCTRINE.md)

binding law: an automation level A_n is CAPPED by the danger tier it is rated for.
 A process may self-move autonomously only up to its permitted L.

Two governing truths:

1. **Loop-risk = act-risk + repetition-risk.** A harmless act repeated a thousand times can become dangerous through fanout, cost, noise, or attention saturation. A loop is never judged by one iteration.
2. **No closure, no automation.** A process that cannot declare its trigger, its limit, its evidence, and its closure does not automate — it may still exist and run manually.

The detailed A0–A5 ladder (trigger, initiator, loop, closure, evidence per level) is the operational reference in docs/AUTOMATION_DOCTRINE.md, mirroring the danger matrix.

Jurisprudence — learning is a process

The Lab learns over time, but **jurisprudence never becomes truth — it becomes tested process compression.** Closed cases may lighten future treatment only through a process:

closed cases

- > projection of the recurring pattern (non-authoritative)
- > hypothesis of a scoped rule (declares scope + counterexamples + supporting case hashes)
- > replay against history (would it have erred?)
- > shadow mode (predicts but does not move the runtime)

- > promotion proportional to risk
- > active rule with a review / sunset / supersession condition

Confidence never becomes authority; it only selects *which rite may be attempted*. Low-risk precedent (L0/L1) may compress quickly; L2 uses shadow/review; L3+ needs explicit grants; **L4/L5 never emerge from precedent alone**. Learning itself registers, is evaluated, carries risk, leaves evidence, and may be doubted or superseded — it enters through the front door.

Governed autonomy

scale + learning over time = governed autonomy

Autonomy is **not** permissionless action. It is **earned compression of repeated judgment**: the Lab handles more, faster, with less direct intervention, because repeated registered processes have taught it which movements can safely become lighter. Stages, roughly:

- A0 manual Dan points; Lab records.
- A1 assisted LLM digests; human approves; Lab registers.
- A2 patterned Lab recognizes repeated low-risk shapes and suggests process.
- A3 fast-path L0/L1 processes move with light or no review.
- A4 governed-auto reversible/low-risk processes execute under executor rules.
- A5 delegated specific higher-risk processes run under explicit grants/budgets/scopes/evl
- A6 institutional standing routines, audits, safe closures with minimal intervention.

The enzyme boundary

The LLM is the **digestive enzyme** of the Lab, not its sovereign organ. The permitted miracle:

language -> structure

The forbidden moves (for the LLM, and by Law 1, for any origin asserting itself):

- structure -> authority (self-ratification)
- authority -> consequence (executing its own proposal, mutating canon, closing its own lo

The LLM may classify, summarize, route, mine, structure, propose, raise attention, and emit candidate Acts (pre-registration). It may not execute code, send externally, mutate canon, or ratify itself. **Safety comes from process classification, not from trusting any origin.**

Risks and hardeners

Risk	Hardener
Reflexivity — the Lab describes itself instead of improving life/work	Operational rent: every canon object must cash out into fields, tests, contracts, projections, attention, or decisions. Enforced by a watcher (below), not willpower.
Digestion loss — LLM translation is lossy and non-deterministic	The Lab does not freeze pure intention; it freezes a content-addressed interpretation with author, source, model, context, doubt, and ratification. Truth is born when the translation is registered, evaluated, and closed by process.
Authority inflation — “the LLM understood” mistaken for “this is true”	Law 1 + the enzyme boundary: movement comes only through process; nothing is true by origin.
Livability failure — elegant but does not make Dan’s days lighter	The Day-8 livability test below.

Operational rent, made self-enforcing. “Every beauty must pay operational rent” is itself a process: a low-tier recurring watcher (e.g. `canon_rent_audit.v1`, danger L1 / automation A2) flags any canon Act that, after a review window, has produced **zero** downstream Acts, tests, projections, contracts, or attention — surfacing it as **attention**, not deleting it. The doctrine against liturgy is enforced by the runtime.

The Day-8 livability test. The Lab is not successful because it is coherent. It is successful when it reduces friction in real life and business operation. Track it as a projection: human-correction rate, time-to-closure, attention backlog. *Build the cathedral — and keep checking that you can still live in it.*

Decision (the laws this record sets)

1. **Process, not source privilege** — origin decides nothing; process decides treatment, risk the rite, consequence the weight.
2. **Risk is consequence** — the L0–L5 danger matrix classifies what an act *does*.
3. **Scale = fixed form + processual classification + proportional risk.**
4. **Automation is continuity (A0–A5), capped by danger (L); loop-risk = act-risk + repetition-risk; no closure, no automation.**
5. **Jurisprudence is learning-as-process** — precedent compresses into scoped, tested, reviewable rules; it never becomes truth; L4/L5 never auto-emerge.

6. **Autonomy is earned compression of repeated judgment** — governed continuity, never permissionless.
7. **The enzyme boundary** — any origin (LLM or human) gains movement only through process; the LLM may reveal structure but may not self-ratify, execute, mutate canon, or close its loop.
8. **Operational rent** — canon that cannot become Act, test, contract, projection, attention, or decision is literature, not runtime priority; a watcher enforces this.

Consequences

- The Lab needs no special metaphysics for LLM vs human input — one process machinery serves both, proportionally.
- Autonomy grows exactly where it is safe to grow (low danger, reversible, well-precedented), and stays governed where consequence is high.
- The anti-cathedral and livability laws are themselves processes, so the doctrine is self-checking rather than aspirational.

Open reconciliations

1. Exact A0–A5 L0–L5 cap table (drafted in `AUTOMATION_DOCTRINE.md`; tune against real cases).
2. The `did` family for jurisprudence rules (e.g. `rule`, `precedent`) and for the rent watcher.
3. How `canon_rent_audit` measures “downstream” (citation graph over `content_hash`).
4. Reconcile Universal Inbox v4/v5 detail into the automation/attention vocabulary (ADR-0003).

Status

`registered` / accepted. Complementary to ADR-0001; supersede only by a later `did=decision` receipt citing this record’s hash — never by edit-in-place.